

Supervised Learning

(Classification: Bayesian Concept Learning)



Bayesian Classifier Concept

- A **Bayesian classifier** is a type of **probabilistic classifier** that applies **Bayesian probability theory** to make predictions or classifications. It models the probability distribution of different classes or categories given the observed data or features.
- There are several types of Bayesian classifiers, each with its own characteristics and applications such as : **Naive Bayes Classifier**.
- The foundation of Bayesian classifiers, including Naive Bayes and other Bayesian models, is **Bayesian probability theory**.

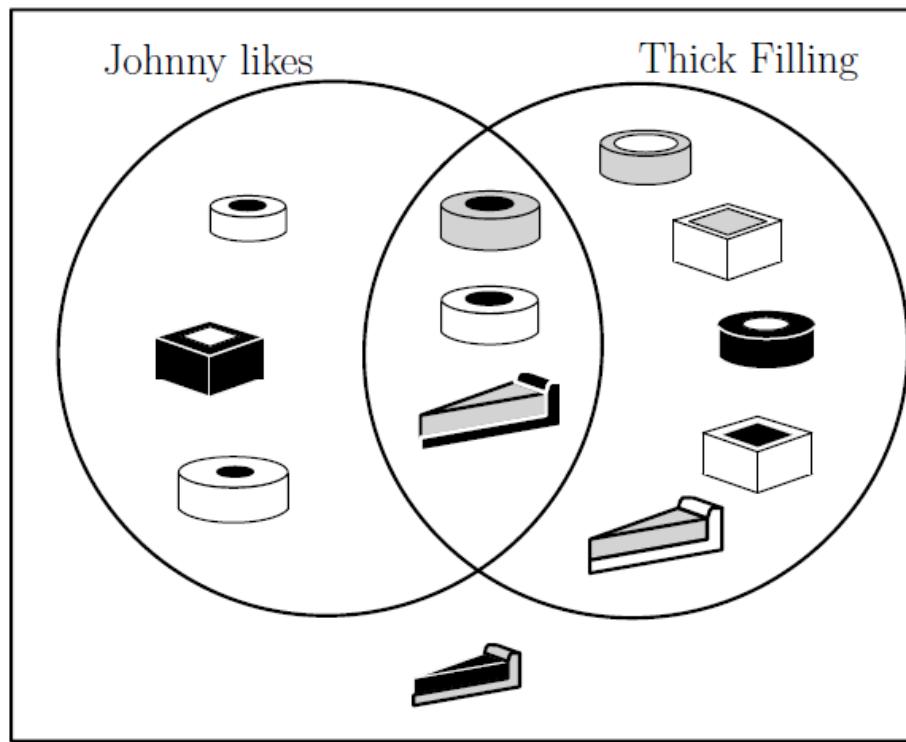
Example of Bayes Theorem

- Given:
 - A doctor knows that meningitis causes stiff neck 50% of the time
 - Prior probability of any patient having meningitis is 1/50,000
 - Prior probability of any patient having stiff neck is 1/20
- If a patient has stiff neck, what's the probability he/she has meningitis?

$$P(M|S) = \frac{P(S|M)P(M)}{P(S)} = \frac{0.5 \times 1/50000}{1/20} = 0.0002$$



Training set in the pies domain



Basic probabilities (revision)

- Relative frequency:

$$P(pos) = \frac{N_{pos}}{N_{all}} = \frac{6}{12} = 0.5$$

- Conditional probability:

$$P(pos | thick) = \frac{N_{pos|thick}}{N_{thick}} = \frac{3}{8} = 0.375$$

- Joint probability:

$$P(pos, thick) = P(pos | thick) \cdot P(thick) = \frac{3}{8} \cdot \frac{8}{12} = \frac{3}{12}$$

Bayesian formula: derivation

- Joint probability is commutative. Therefore:

$$P(pos, thick) = P(pos \mid thick) \cdot P(thick) = P(thick \mid pos) \cdot P(pos)$$

- From here, the Bayesian formula is obtained:

$$P(pos \mid thick) = \frac{P(thick \mid pos) \cdot P(pos)}{P(thick)}$$

Bayesian formula used for classification

- Probability that a pie with *thick-filling* is positive:

$$P(pos \mid thick) = \frac{P(thick \mid pos) \cdot P(pos)}{P(thick)}$$

- Probability that a pie with *thick-filling* is negative:

$$P(neg \mid thick) = \frac{P(thick \mid neg) \cdot P(neg)}{P(thick)}$$

- Choose the label with the higher probability.

The “naïve Bayes” approach

- Probability of class c_j : $P(c_j \mid \mathbf{x}) = \frac{P(\mathbf{x} \mid c_j) \cdot P(c_j)}{P(\mathbf{x})}$
- Naïve assumption: mutually independent attributes, x_i

$$P(\mathbf{x} \mid c_j) = \prod_{i=1}^n P(x_i \mid c_j)$$

- The whole Bayes formula

$$P(c_j \mid \mathbf{x}) = \frac{P(c_j) \cdot \prod_{i=1}^n P(x_i \mid c_j)}{P(\mathbf{x})}$$

“Naïve Bayes”: one practical comment

- Probability of class c_j :
$$P(c_j | \mathbf{x}) = \frac{P(\mathbf{x} | c_j) \cdot P(c_j)}{P(\mathbf{x})}$$
- $P(\mathbf{x})$ is the same for each class, c_j .
- Therefore: choose c_j with the greatest numerator, $P(\mathbf{x}|c_j)P(c_j)$.



“Naïve Bayes” (algorithm)

The example to be classified is described by $\mathbf{x} = (x_1, \dots, x_n)$.

1. For each x_i , and for each class c_j , calculate the conditional probability, $P(x_i|c_j)$, as the relative frequency of x_i among those training examples that belong to c_j .
2. For each class, c_j , carry out the following two steps:
 - i) estimate $P(c_j)$ as the relative frequency of this class in the training set;
 - ii) calculate the conditional probability, $P(\mathbf{x}|c_j)$, using the “naive” assumption of mutually independent attributes:

$$P(\mathbf{x}|c_j) = \prod_{i=1}^n P(x_i|c_j)$$

-
3. Choose the class with the highest value of $P(c_j) \cdot \prod_{i=1}^n P(x_i|c_j)$.
-

Bayes' Theorem

- Bayes' theorem is a fundamental principle in probability theory and statistics that relates conditional probabilities.
- It describes how to update the probability of a hypothesis (an event or statement) based on new evidence or observations.
- The theorem can be expressed mathematically as:

$$P(H|E) = \frac{P(E|H) \cdot P(H)}{P(E)}$$

Where:

- $P(H|E)$ is the posterior probability of hypothesis H being true given evidence E.
- $P(E|H)$ is the probability of observing evidence E given that hypothesis H is true (the likelihood).
- $P(H)$ is the prior probability of hypothesis H being true before considering the evidence.
- $P(E)$ is the probability of observing evidence E, which can be calculated as a marginal probability by considering all possible hypotheses.

Naïve Bayes (example)

Ex.	Crust size	shape	Filling size	class
e1	big	circle	small	pos
e2	small	circle	small	pos
e3	big	sq.	big	pos
e4	small	sq.	big	pos
e5	big	circle	big	pos
e6	big	sq.	small	neg
e7	big	tri.	small	neg
e8	small	sq.	big	neg

$$\begin{aligned}P(\text{neg}) &= 3/8; \\ P(\text{pos}) &= 5/8\end{aligned}$$

$\mathbf{x}=(\text{small}, \text{square}, \text{small})$

- $P(\text{crust}=\text{small} | \text{pos}) = 2/5$
- $P(\text{shape}=\text{square} | \text{pos}) = 2/5$
- $P(\text{filling}=\text{small} | \text{pos}) = 2/5$
- $P(\mathbf{x} | \text{pos}) = \frac{2}{5} \times \frac{2}{5} \times \frac{2}{5} = \frac{8}{125}$
- $P(\text{crust}=\text{small} | \text{neg}) = 1/3$
- $P(\text{shape}=\text{square} | \text{neg}) = 2/3$
- $P(\text{filling}=\text{small} | \text{neg}) = 2/3$
- $P(\mathbf{x} | \text{neg}) = \frac{1}{3} \times \frac{2}{3} \times \frac{2}{3} = \frac{4}{27}$
- $P(\mathbf{x} | \text{pos}) P(\text{pos}) = \frac{8}{125} \times \frac{5}{8} = 0.040$
- $P(\mathbf{x} | \text{neg}) P(\text{neg}) = \frac{4}{27} \times \frac{3}{8} = 0.056$
- *Verdict: neg because $0.056 > 0.040$*



Naïve Bayes (another example)

- $\mathbf{x} = [\text{square}, \text{thick}, \text{gray}, \text{thin}, \text{white}]$

$P(\text{shape}=\text{square} \text{pos})$	= 1/6	$P(\text{shape}=\text{square} \text{neg})$	= 2/6
$P(\text{crust-size}=\text{thick} \text{pos})$	= 5/6	$P(\text{crust-size}=\text{thick} \text{neg})$	= 5/6
$P(\text{crust-shade}=\text{gray} \text{pos})$	= 1/6	$P(\text{crust-shade}=\text{gray} \text{neg})$	= 2/6
$P(\text{filling-size}=\text{thin} \text{pos})$	= 3/6	$P(\text{filling-size}=\text{thin} \text{neg})$	= 1/6
$P(\text{filling-shade}=\text{white} \text{pos})$	= 1/6	$P(\text{filling-shade}=\text{white} \text{neg})$	= 2/6

We see that $P(\text{pos}) = P(\text{neg}) = 0.5$ and also:

$$P(\mathbf{x}|\text{pos}) = \prod_{i=1}^n P(x_i|\text{pos}) = \frac{1}{6} \cdot \frac{5}{6} \cdot \frac{1}{6} \cdot \frac{3}{6} \cdot \frac{1}{6} = \frac{15}{6^5}$$

$$P(\mathbf{x}|\text{neg}) = \prod_{i=1}^n P(x_i|\text{neg}) = \frac{2}{6} \cdot \frac{5}{6} \cdot \frac{2}{6} \cdot \frac{1}{6} \cdot \frac{2}{6} = \frac{40}{6^5}$$

Since $P(\mathbf{x}|\text{pos}) \cdot P(\text{pos}) < P(\mathbf{x}|\text{neg}) \cdot P(\text{neg})$, we label \mathbf{x} with the negative class.

Naïve Bayes Classifier with Python

Imagine you are a data scientist working on a project to classify iris flowers into different species based on their attributes. You have a dataset with labeled examples for training and testing purposes. Your goal is to develop a machine learning model using the Gaussian Naive Bayes algorithm to classify iris flowers accurately.

Note : The iris dataset consists of continuous numerical features, including sepal length, sepal width, petal length, and petal width. Gaussian Naive Bayes is well-suited for datasets with continuous features because it assumes that these features follow a Gaussian (normal) distribution.

To import the GaussianNB class from Scikit-Learn's `naive_bayes` module, you can use the following import statement in Python:

```
from sklearn.naive_bayes import GaussianNB
```

This import statement allows you to create an instance of the GaussianNB classifier as follows:

```
nb = GaussianNB()
```

Upload Iris Dataset.

These steps will load the Iris dataset, create a DataFrame for it, perform a train/test split, and print the shapes of the resulting datasets.

```
# Import necessary libraries
from sklearn import datasets
import pandas as pd
import sklearn.model_selection as skms
# Load the Iris dataset
iris = datasets.load_iris()
# Create a DataFrame for the dataset
iris_df = pd.DataFrame(iris.data, columns=iris.feature_names)
iris_df['target'] = iris.target
# Simple train/test split of the dataset
(iris_train_ftrs, iris_test_ftrs, iris_train_tgt, iris_test_tgt) = skms.train_test_split(iris.data,
iris.target, test_size=.25)
# Print the shapes of the training and testing sets
print("Train features shape:", iris_train_ftrs.shape)
print("Test features shape:", iris_test_ftrs.shape)
```

Naïve Bayes Classifier with Python

```
nb = naive_bayes.GaussianNB()
```



In this line, you create an instance of the Gaussian Naive Bayes classifier by calling the GaussianNB() constructor from the naive_bayes module of scikit-learn. This classifier assumes that the features (attributes) are continuous and normally distributed.

```
fit = nb.fit(iris_train_ftrs, iris_train_tgt)
```



Here, you use the fit method to train (fit) the Gaussian Naive Bayes classifier on your training data. iris_train_ftrs represents the training features (attribute values) from your dataset, and iris_train_tgt represents the corresponding target labels (class labels) for the training examples. This step involves estimating the class-conditional probabilities and other parameters required for classification..

```
preds = fit.predict(iris_test_ftrs)
```



Here, you use the fit method to train (fit) the Gaussian Naive Bayes classifier on your training data. iris_train_ftrs represents the training features (attribute values) from your dataset, and iris_train_tgt represents the corresponding target labels (class labels) for the training examples. This step involves estimating the class-conditional probabilities and other parameters required for classification..

Naïve Bayes (Summary)

- Robust to isolated noise points
- Handle missing values by ignoring the instance during probability estimate calculations
- Robust to irrelevant attributes